



# CSIS 429 Operating Systems

## Lecture 2: CPU Mechanisms

September 9<sup>th</sup> 2020

# Textbook chapters

Read “Processes” and “Process API”

Intro	Virtualization		Concurrency	Persistence	Appendices
<a href="#">Preface</a>	<a href="#">3 Dialogue</a>	<a href="#">12 Dialogue</a>	<a href="#">25 Dialogue</a>	<a href="#">35 Dialogue</a>	<a href="#">Dialogue</a>
<a href="#">TOC</a>	<a href="#">4 Processes</a>	<a href="#">13 Address Spaces</a>	<a href="#">26 Concurrency and Threads code</a>	<a href="#">36 I/O Devices</a>	<a href="#">Virtual Machines</a>
<a href="#">1 Dialogue</a>	<a href="#">5 Process API code</a>	<a href="#">14 Memory API</a>	<a href="#">27 Thread API</a>	<a href="#">37 Hard Disk Drives</a>	<a href="#">Dialogue</a>
<a href="#">2 Introduction code</a>	<a href="#">6 Direct Execution</a>	<a href="#">15 Address Translation</a>	<a href="#">28 Locks</a>	<a href="#">38 Redundant Disk Arrays (RAID)</a>	<a href="#">Monitors</a>
	<a href="#">7 CPU Scheduling</a>	<a href="#">16 Segmentation</a>	<a href="#">29 Locked Data Structures</a>	<a href="#">39 Files and Directories</a>	<a href="#">Dialogue</a>
	<a href="#">8 Multi-level Feedback</a>	<a href="#">17 Free Space Management</a>	<a href="#">30 Condition Variables</a>	<a href="#">40 File System Implementation</a>	<a href="#">Lab Tutorial</a>
	<a href="#">9 Lottery Scheduling code</a>	<a href="#">18 Introduction to Paging</a>	<a href="#">31 Semaphores</a>	<a href="#">41 Fast File System (FFS)</a>	<a href="#">Systems Labs</a>
	<a href="#">10 Multi-CPU Scheduling</a>	<a href="#">19 Translation Lookaside Buffers</a>	<a href="#">32 Concurrency Bugs</a>	<a href="#">42 FSCK and Journaling</a>	<a href="#">xv6 Labs</a>
	<a href="#">11 Summary</a>	<a href="#">20 Advanced Page Tables</a>	<a href="#">33 Event-based Concurrency</a>	<a href="#">43 Log-structured File System (LFS)</a>	<a href="#">Flash-based SSDs</a>
		<a href="#">21 Swapping: Mechanisms</a>	<a href="#">34 Summary</a>	<a href="#">44 Data Integrity and Protection</a>	
		<a href="#">22 Swapping: Policies</a>		<a href="#">45 Summary</a>	
		<a href="#">23 Case Study: VAX/VMS</a>		<a href="#">46 Dialogue</a>	
		<a href="#">24 Summary</a>		<a href="#">47 Distributed Systems</a>	
				<a href="#">48 Network File System (NFS)</a>	
				<a href="#">49 Andrew File System (AFS)</a>	
				<a href="#">50 Summary</a>	

# Understanding Operating Systems

- Operating system
  - Software that helps other programs control computer hardware and interact with users

## Application

- Software program that provides service for computer user
- Cannot act without “permission” from operating system

# Operating System Components

Major operating system components:

- Kernel
- Device drivers
- Shell
- Utility programs
- Graphical user interface (GUI)

# Operating System Functions

- Abstraction: provide system calls and libraries for using resources
- Manage system resources

# “Crux of the problem”

The main idea of the entire course

**How** does an OS virtualize:

- ✓ the CPU
  - ✓ Memory
  - ✓ Disk
- } Resources

**efficiently?**

# Operating System: Abstraction

Provide system calls and libraries for using resources

- Resources are things like CPU, memory, disk, etc.
- How does an OS “abstract” resources?
  - CPU: abstracted as processes or threads
  - Memory: process address space
  - Disk: files
- Why use abstraction?
  - Make the system easy for developers to use
  - Can make different devices look the same – e.g. files

# Operating System: Managing Resources

- Why should an OS manage resources?
  - Security: protect applications from each other
  - Fairness
    - users should be able to access & use the system
    - Disk: files
  - Efficiency
    - Allow efficient (cost, time, energy) access to h/w



# Operating System == Virtual Machine

- An OS provides APIs – a standard library for applications

Other names for Oses:

- Supervisor
- Master Control Program

# The UNIX Operating System

## UNIX

- An operating system
- Originally created at AT&T Bell Labs in early 1970s
- Designed to control networked computers that were shared by many users
- Features and low cost of Linux effectively driving UNIX out of market

# Three Easy Pieces

- The three easy pieces are the three main themes in operating systems:
  - 1) Virtualization – make each application believe it has the resource to itself.
  - 2) Concurrency – handle interactions between processes running at the same time
  - 3) Persistence – data should stick around, beyond the lifetime of an application run

# Hardware Resource: CPU

What did you learn about CPUs from CSIS 355?

What does a CPU do?

# Hardware Resource: CPU

CPUs perform a very simple set of tasks:

- 1) Fetch an instruction from memory
  - 2) Decode the instruction
  - 3) Execute
- Repeat!

Very quickly – at about a billion instructions/second

# Programs

For all their complexity, programs also perform a simple set of tasks:

- 1) Load code and data into memory
- 2) Fetch an instruction from memory
- 3) Execute

Repeat steps 2 & 3



# Programs

For all their complexity, programs also perform a simple set of tasks:

- 1) Load code and data into memory
- 2) Fetch an instruction from memory
- 3) Execute

Repeat steps 2 & 3



Notice:

A program does the same thing as the CPU

The OS “virtualizes” the CPU so that each program thinks it is the only thing running

# Virtualization

Let's take a look at

- `cpu.c`

Compile using `gcc cpu.c -lpthread`

Can run multiple instances:

`./cpu P1 & ./cpu P2 & ./cpu P3`



# Virtualizing the CPU

○ The main way to virtualize the CPU is to use the "process" abstraction:

A process is an execution stream within the context of an execution state.

Execution stream == stream of instructions executing, thread of control

Execution state == everything that affects the instruction stream: CPU registers, heap, stack, open files, other data.

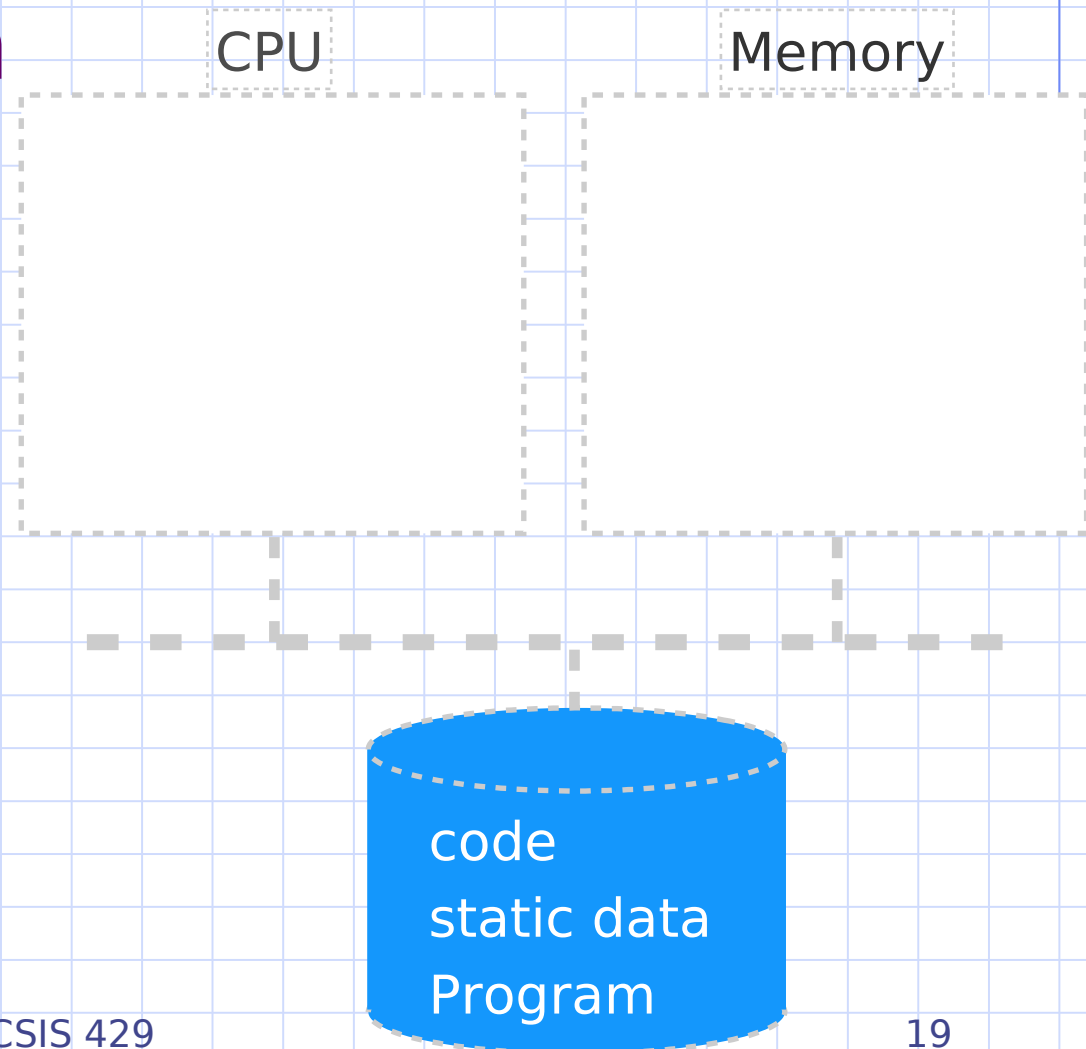
# Processes vs. Programs

Program: static code, static data

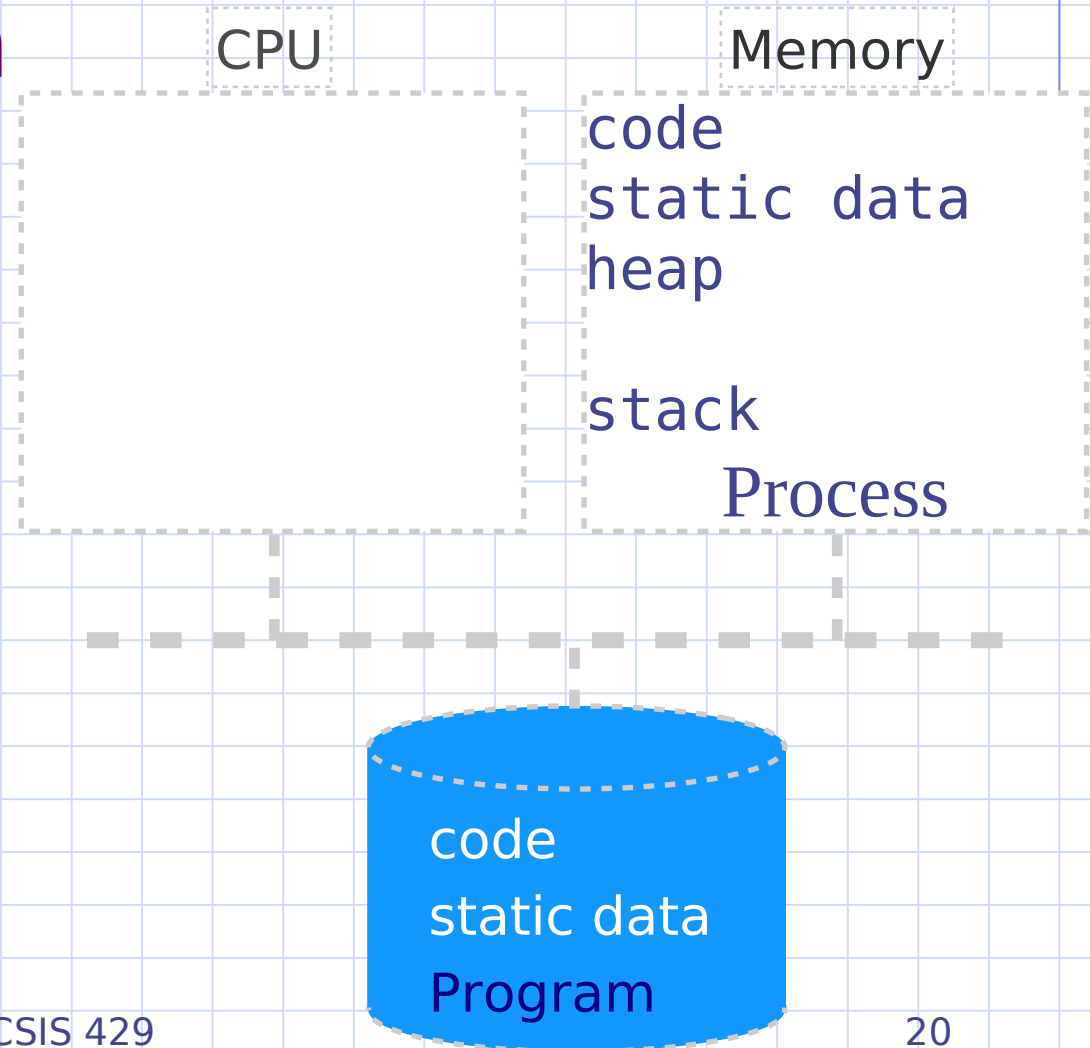
Process: code undergoing execution, changing data

We can run multiple instances (processes) of the same program

# Processes Creation



# Processes Creation



In Chapter 4 (Processes)  
read section 4.3 (pages 4-  
5) on **Process Creation**  
for more details

# Virtualizing Memory

Let's take a look at

- `mem.c`

Compile using `gcc`  
`./mem X`

Look at addresses printed by `mem`  
ASLR – can be turned off in **`gdb`**

```
(gdb) set disable-randomization off  
(gdb) show disable-randomization
```

*Based on what you learned in CSIS 248 (Operating System Programming):*

When we run multiple instances of `mem`, are we accessing the same physical RAM location?

# Virtualizing Memory

When we run multiple instances of `mem`, are we accessing the same physical RAM location?

OS provides a “process address space”

Also called “virtual address space” or “virtual memory”

→ OS “maps” these to physical RAM locations

# Direct vs. Controlled program execution

Direct: allow a user process to run directly on hardware - “DOS” model. Can break:

- Security – can change files, settings, etc.
- Fairness – can use CPU forever
- Efficiency – may busy-wait for slow I/O

Solution: OS and hardware maintain some control over what processes can do.

# How can we control program execution

- Set up two levels: user (restricted) mode and kernel (unrestricted) mode
  - ➔ User processes run in user mode
  - ➔ OS runs in kernel mode

## How can we enforce security?

- Processes have to use special “system calls” to access valuable resources
- System calls are checked by kernel



# OS tracks program execution

- Each process is in one state at any single time:

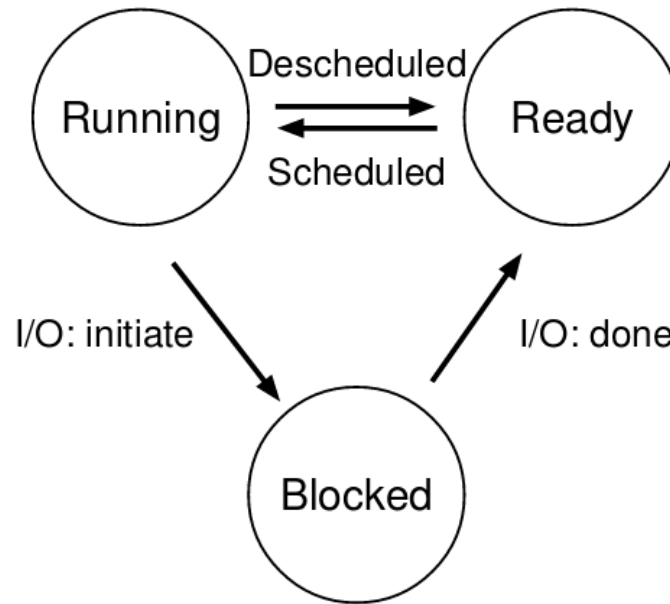



Figure 4.2: Process: State Transitions

# System Calls

Process A wants to call `sys_read()` to access disk

A is in user space and can only see its own memory: other processes and kernel are “hidden” → no way to call `sys_read` directly.

Instead, it calls the `read()` library function



# System Calls

Process A has to call the read() library function

read():

```
movl $6, %eax; %syscall table index  
int $64; %trap table index
```

RAM




# System Calls

Once process A executes the interrupt, we switch to Kernel Mode  
→ executing as kernel, we can do anything!

read():


```
movl $6, %eax; %syscall table index  
int $64; %trap table index
```



# System Calls

Once process A executes the interrupt, we switch to Kernel Mode  
→ executing as kernel, we can do anything!

Kernel can access user memory to fill the buf buffer space when done reading from disk.



# What are user processes not allowed to do?

- User processes are not allowed to
  - Access memory beyond process address space
  - Access disk
  - Execute privileged (ring 0) x86 instructions

# Privileged (ring 0) x86 instructions



LGDT	Loads an address of a GDT into GDTR
LLDT	Loads an address of a LDT into LDTR
LTR	Loads a Task Register into TR
MOV Control Register	Copy data and store in Control Registers
LMSW	Load a new Machine Status WORD
CLTS	Clear Task Switch Flag in Control Register CR0
MOV Debug Register	Copy data and store in debug registers
INVD	Invalidate Cache without writeback
INVLPG	Invalidate TLB Entry
WBINVD	Invalidate Cache with writeback
HLT	Halt Processor
RDMSR	Read Model Specific Registers (MSR)
WRMSR	Write Model Specific Registers (MSR)
RDPMC	Read Performance Monitoring Counter
RDTSC	Read time Stamp Counter

# How does an OS control access to CPU?

• In order to allow an OS to switch between processes, the OS has to be able to “stop” one process, save its “execution state” and switch to another process:

- Run process A for its “time slice”
- Stop process A execution and save context of A
- Load context of process B
- Start process B execution



# OS controlling access to CPU

Operating System

Hardware

Program  
Process A

timer interrupt  
**context-switch**: save regs(A) to  
kernel-stack(A)  
move to kernel mode  
jump to trap handler

Handle the trap  
Call **switch()** routine  
save regs(A) to proc-struct(A)  
restore regs(B) from proc-  
struct(B)  
switch to k-stack(B)  
return-from-trap (into B)

**Context-switch**: restore regs(B)  
from k-stack(B)  
move to user mode  
jump to B's IP

Process B

# CPU Virtualization

- The goal of CPU virtualization is to run  $N$  processes on  $M$  CPUs.  $N \gg M$

Abstraction used: “process”

Each process “thinks” it has access to the entire machine.

# CPU Virtualization

- To support the “process” abstraction, we need a lot of support in the kernel and CPU.

Each process “thinks” it has access to the entire machine.

# Concurrency

Let's take a look at

- `threadv0.c`
- `threadv1.c`

Compile using `gcc`  
May need `-lpthread`